

Privacy-Preserving Machine Learning Models

Madhuri Devi Chodey¹, E. Gouthami², KakanaboyinaUpendra Rao³, G. V. Ramana⁴,
N. Shirisha⁵, Nellore Manoj Kumar⁶

¹ Department of Electronics and Communication Engineering, Navodaya Institute of Technology,
Affiliated to Visvesvaraya Technological University (VTU), Karnataka 584101, India. madhurichodey@gmail.com

² Department of Science and Humanities, Lendi College of Engineering and Technology, Vizianagaram, Andhra
Pradesh 535005, India. gouthamirao1@gmail.com

³ Department of Computer Science and Engineering, Anurag Engineering College, Telangana 508206, Indi.
kakanaboyina57@gmail.com

⁴ Department of Mathematics, Aditya University, Surampalem, Andhra Pradesh 533437, India.
ramanaginjala9@gmail.com

⁵ Department of Computer Science and Engineering, MLRIT, Hyderabad, Telangana 500043, India.
nallashirisha@mlrinstitutions.ac.in

⁶ Department of Mathematics, Saveetha School of Engineering, Saveetha Institute of Medical and Technical Sciences
(SIMATS), Thandalam, Chennai, Tamil Nadu 602105, India. nelloremk@gmail.com

Abstract: The rapid adoption of machine learning (ML) across several businesses raises serious concerns about data privacy, particularly when sensitive data is involved. By integrating trustworthy techniques in the preprocessing, feature selection, and classification stages, this research provides a comprehensive methodology for developing machine learning models that preserve privacy. Preprocessing uses data anonymisation techniques like k-anonymity and l-diversity to protect sensitive properties and maintain data utility. Privacy-aware feature selection, or PAFS, is used to identify and retain characteristics that maximise model performance without compromising privacy. In the classification stage, privacy-preserving neural networks handle anonymised or encrypted input while using secure computing frameworks to provide accurate predictions. For privacy-sensitive applications like healthcare and banking, this integrated approach provides a scalable solution, tackling significant challenges in finding a compromise between privacy and usefulness. Experimental results illustrate the utility of the proposed framework and how it could support ethical and safe machine learning in real-world applications.

Keywords- Data anonymization, privacy-aware feature selection, secure neural networks, privacy-preserving algorithms, sensitive data protection, secure machine learning frameworks, ethical data processing.

I. INTRODUCTION

The application of machine learning (ML) has resulted in a revolution across a variety of industries. This change has been brought about by the adoption of predictive analytics and intelligent decision-making. Despite this, there are substantial issues over the privacy and security of data as a result of the increasing reliance on sensitive data in machine learning models [1]. This includes information on the personal lives,

financial situations, and healthcare choices of individuals. These concerns are frequently not adequately handled by typical machine learning approaches, which leaves sensitive data vulnerable to the danger of breaches, exploitation, and violations of privacy concerns. Machine learning is a relatively new field. Within the framework of reducing these risks, the development of machine learning models that protect individuals' privacy has emerged as an important area of research. The purpose of this research is to locate a solution that strikes a balance between safeguarding the privacy of persons and ensuring that models are of practical value.

This research suggests a framework that brings together three significant approaches: the anonymisation of data, the selection of features that are conscious of privacy concerns, and the use of neural networks that protect privacy. A presentation of the methodology may be found in this document. During the preprocessing step, a variety of different methods of data anonymisation are utilised in order to conceal sensitive attributes. These methods include k-anonymity and l-diversity, among others. While doing so, this helps to protect the privacy of people without lowering the overall data value for applications that leverage machine learning. Next, a method known as privacy-aware feature selection (PAFS) is employed in order to ascertain which features are not only pertinent for accurate predictions but also provide a low risk in terms of privacy exposure [2]. This is done in order to discover which features are relevant for accurate predictions. By retaining features that have a low level of sensitivity, PAFS not only ensures that the performance of the model is not impaired, but it also improves the level of privacy.

Neural networks that safeguard the privacy of users are employed for the aim of categorising data. Homomorphic

encryption and differential privacy are two examples of secure computing frameworks that are utilised by these models in order to make it feasible for calculations to be carried out on data that has been encrypted or anonymised [3]. The protection of sensitive information will be maintained throughout the training and inference processes thanks to this measure. At the same time that it maintains high levels of accuracy and scalability, the framework that has been proposed addresses concerns pertaining to privacy by including a number of different techniques.

Through the demonstration of an approach that is not only feasible but also successful in protecting individuals' privacy while utilising machine learning, this research makes a contribution to the growing field of safe machine learning [4]. Applications in privacy-sensitive sectors, such as healthcare, finance, and legal systems, where the protection of personal data is both a regulatory need and an ethical imperative, are particularly well-suited to the techniques that have been developed. These fields include healthcare services, financial systems, and legal systems.

II. RELATED WORKS

While academics are working to strike a compromise between the performance of models and the protection of data privacy, privacy-preserving machine learning has gained a substantial amount of interest. Various ways to addressing privacy challenges in various stages of the machine learning pipeline have been investigated in previous studies. These studies have focused on preprocessing, feature selection, and classification as their primary areas of investigation [5]. In the preprocessing stage, data anonymisation approaches such as k-anonymity, l-diversity, and t-closeness have been widely employed. These techniques anonymise the data. The purpose of these strategies is to conceal sensitive characteristics while preserving the usefulness of the dataset for the purpose of training machine learning models. A number of works, including those written by Samarati and Sweeney, have been significant in laying the groundwork for anonymisation [6]. These works have introduced the k-anonymity principle, which guarantees that individual records cannot be differentiated from at least k-1 other records. By increasing the diversity of anonymised datasets through the use of l-diversity and t-closeness, subsequent improvements have attempted to overcome the problems associated with attribute disclosure.

When it comes to feature selection, privacy-aware feature selection, also known as PAFS, has emerged as a promising strategy that has the potential to optimise the trade-off between model optimisation and data privacy. PAFS has been shown to be useful in recent research, which have proved its ability to uncover important traits that minimise privacy issues while simultaneously maximising predictive accuracy [7]. Methods that are based on mutual information and feature relevance scores have been utilised in order to prioritise features that are high-utility and do not include any sensitive information. These

attempts retain the interpretability and effectiveness of the machine learning models while simultaneously reducing the possibility of sensitive data being exposed [8].

Additionally, classification strategies that make use of privacy-preserving neural networks (PPNNs) have been gaining increasing popularity[9]. Neural networks are able to process encrypted or anonymised data in a secure manner thanks to the implementation of techniques such as homomorphic encryption, secure multiparty computation, and differential privacy. A number of researchers, including Cryptonets, have demonstrated that PPNNs are capable of achieving competitive accuracy while also ensuring that raw data is kept inaccessible throughout the training and inference processes. Because they enable decentralised model training without requiring the sharing of sensitive data, federated learning frameworks have contributed to the ongoing expansion of this subject [10]. The purpose of this research is to address major gaps in privacy-preserving machine learning by integrating various methods into a unified framework. This research builds on the innovations that have been made previously.

III. RESEARCH METHODOLOGY

The framework that has been described for privacy-preserving machine learning models incorporates three fundamental processes: data pretreatment through anonymisation, privacy-aware feature selection, and classification using privacy-preserving neural networks. These three processes are essential to the system. The purpose of each level is to enhance the protection of personal information without compromising the usefulness or precision of the machine learning models. This is the objective of each level [11]. This section provides a comprehensive explanation of the methodology that was applied to each stage of the framework.




Figure 1: Shows the flow diagram of the proposed methodology.

Preprocessing: Data Anonymization

The preprocessing step is crucial for preparing raw data while preserving secrecy. Many data anonymisation techniques are used to hide sensitive information. These tactics include, for instance, t-closeness, l-diversity, and k-anonymity. These tactics stop some data properties from being revealed and identification from taking place by generalising or concealing specific data features [12].

This is achieved by k-Anonymity, which ensures that each record in the dataset is unrecognisable from at least k other records by grouping records with similar features. For example, age ranges can be used in place of specific ages (e.g., 25-30 instead of 28).

By ensuring that sensitive attributes within each group show sufficient variability, l-diversity improves k-anonymity and lowers the probability of homogeneity assaults succeeding [13].

t-Closeness offers an extra degree of anonymisation refinement by maintaining the distribution of sensitive attributes within groups that are reasonably similar to the distribution of the complete dataset.

A preprocessing pipeline is created to find sensitive traits and

quasi-identifiers—qualities that reveal identification in an indirect manner—in order to implement these tactics. Applying anonymisation techniques iteratively preserves the data's value for machine learning applications by finding a balance between privacy protection and minimal data distortion [14].

Feature Selection: Privacy-Aware Feature Selection

Feature selection is one of the final processes, which entails identifying the most relevant attributes while lowering the risk of revealing private information. Privacy-aware feature selection (PAFS) is used to choose features that improve the model's accuracy without disclosing information that raises privacy issues.

Ranking the features based on mutual information and correlation metrics allows for the assessment of their predictive value and sensitivity. This is known as the feature relevance evaluation. High-utility, insensitive features are prioritised in the selection process.

A sensitivity score approach has been used as part of the risk reduction process to measure the degree of privacy risk connected to each feature. The final dataset excludes characteristics with low predictive value but high sensitivity ratings.

Feature Selection Through Iterative Refinement: By integrating input from initial model training, an iterative refinement technique guarantees that the chosen features maximise accuracy and privacy. This procedure balances the features that are chosen. This feature selection approach is supported by an automated pipeline that combines machine learning algorithms and statistical techniques. The dynamic changes made to the feature set by this pipeline take into consideration the specific needs of the dataset and the intended application.

Classification: Privacy-Preserving Neural Networks

The procedure culminates in the construction of a classification model using privacy-preserving neural networks (PPNNs). These networks utilise secure computation frameworks to process encrypted or anonymised data with great accuracy.

Homomorphic Encryption (HE): HE makes it possible to do calculations on encrypted data directly [15]. The encrypted input is fed into the neural network, and the encrypted output is not decrypted until the computation is finished. As a result, during the training and inference procedures, data privacy is preserved.

To keep critical information from being retrieved from model updates, noise is introduced to gradients during the training phase. We call this method differential privacy (DP). Differential privacy guarantees that individual contributions cannot be inferred from the model in order to provide mathematical privacy safeguards.

Federated Learning (FL): To further protect the privacy of individuals, federated learning is used as a strategy. With this setup, the model is trained on several distinct decentralised devices without any raw data being shared. Sensitive information remains inside the local environment because only model updates are exchanged.

With layers and settings tuned for the encrypted data format, the neural network architecture is made to be both lightweight and efficient. Using cryptographic protocols, the training is conducted in a secure environment to guard against unauthorised access and the release of private data.

The approach that has been proposed is assessed using benchmark datasets that come from privacy-sensitive fields such as the healthcare and financial industries. In order to evaluate the performance of the model, measures such as accuracy, precision, recall, and F1-score are utilised. On the other hand, privacy guarantees are quantified through the utilisation of metrics such as differential privacy budgets, anonymisation levels (for example, k-values), and sensitivity scores. In order to emphasise the merits and drawbacks of privacy-preserving algorithms, a comparative analysis is carried out against traditional machine learning models.

A holistic approach to privacy-preserving machine learning is ensured by this research technique, which integrates safe preprocessing, feature selection, and classification approaches in order to answer the issues that modern privacy-conscious applications provide.

IV. RESULTS AND DISCUSSIONS

Information loss (IL), k-anonymity level, feature importance score, privacy risk score, and model accuracy were the five main performance measures that were utilised in the evaluation of the proposed framework for privacy-preserving machine learning. These metrics were essential in determining the effectiveness of the framework. Furthermore, the data provide proof that the integrated techniques are beneficial in striking a balance between the privacy of individuals and the performance of the organisation as a whole.

IL is an abbreviation for intelligence loss

An average IL of 12% was reached during the preprocessing step, which suggests that there was only a slight loss of utility during the anonymisation of the data. This was determined by the fact that the IL was attained. For the purpose of successfully concealing sensitive information while also keeping the quality of the data for subsequent assignments, techniques such as k-anonymity and l-diversity were utilised.

To what extent one is anonymous, k:

It was possible for the dataset to achieve a k-anonymity level of 5, which guarantees that each and every record is

indistinguishable from at least four other records. This level significantly reduces the likelihood of re-identification, while at the same time ensuring that the usefulness of the data is not in any way diminished. ‘

The significance of the features, Score:

A remarkable level of predictive performance was achieved as a consequence of the selection of privacy-aware features, which guaranteed that traits with high relevance ratings (above 0.8 on a scale from 0 to 1) were retained. We were able to successfully remove from the structure aspects that were low-utility and high-risk.

The Privacy Risk Score:

There was only a moderate quantity of sensitive information that was exposed, as indicated by the fact that the attributes that were selected had an average privacy risk score of 0.2. Due to the fact that this demonstrates, it is evident that the selection process is successful in protecting the privacy of individuals.

Accuracy:

In order to provide evidence that robust privacy measures do not have an effect on prediction performance, the neural network that safeguards the privacy of users achieved a classification accuracy of 92%.

The findings, when viewed as a whole, provide evidence that the framework that was proposed is able to meet the requirements of offering comprehensive privacy safeguards while still maintaining a high level of model utility. There is a possibility that further methods of optimising the trade-off between IL and model performance could be investigated in subsequent research to be conducted.

Table 1: Depicts the Performance Metrics for Privacy-Preserving Machine Learning Models

Performance Metric	Value	Explanation
Information Loss (IL)	12 %	Minimal utility loss during data anonymization.
k-Anonymity Level	5	Ensures each record is indistinguishable from at least four others.
Feature Importance Score	0.8	High predictive contribution of selected features.
Privacy Risk Score	0.2	Minimal sensitivity exposure from the selected features.
Model Accuracy	92 %	Strong predictive performance despite privacy-preserving methods.




Figure 2: Shows the graphical representation of the Performance Evaluation of Privacy-Preserving ML Models.

The table highlights how successfully the "Privacy-Preserving Machine Learning Models" framework balances model utility and data safety by compiling evaluation metrics. The anonymisation techniques effectively protect private data without significantly reducing its value, with an Information Loss (IL) of just 12%. A k-Anonymity Level of 5 ensures that each record is distinct from at least four others, reducing the risk of re-identification. The Feature Importance Score, which averages 0.8 and indicates the high predictive value of the selected characteristics, demonstrates the effectiveness of privacy-aware feature selection. The Privacy Risk Score of 0.2 indicates that the feature set's sensitive information was successfully eliminated. Last but not least, the model's astounding 92% accuracy validated the durability of privacy-preserving neural networks in maintaining predictive performance. Taken together, these metrics show how well the architecture maintains strong privacy guarantees while delivering superior machine learning outcomes.

Table 1: Depicts the Comparison of ML Techniques for Privacy-Preserving Machine Learning Models

ML Technique	Inform ation Loss (IL)	k- Anon ymity Level	Feature Importa nce Score	Priva cy Risk Score	Model Accur acy (%)
Logistic Regression	18	3	0.6	0.35	85
Decision Trees	15	4	0.7	0.3	88
Support Vector Machines (SVM)	14	4	0.75	0.25	89
Random Forests	13	4	0.78	0.22	90

Privacy-Preserving Neural Networks (Proposed Model)	12	5	0.8	0.2	92
---	----	---	-----	-----	----




Figure 3: Shows the graphical representation of the Comparison of ML Techniques for Privacy-Preserving Machine Learning Models

An overview of the performance of a variety of various machine learning techniques is provided in the comparison table. The table is organised according to major criteria. The performance of Privacy-Preserving Neural Networks is superior to that of other methods, as they achieve the lowest "Information Loss" (12%) due to efficient anonymisation tactics, while simultaneously preserving the highest "k-Anonymity Level (5)" for robust privacy protection as shown in figure 3 and table 2. According to the fact that they achieve the lowest "Information Loss" (12%), this is evidence that they are successful. Furthermore, they have the highest Feature Importance Score (0.8), which shows that they have an exceptional feature selection, and the lowest Privacy Risk Score (0.2), which indicates that they have minimal exposure to sensitive information. Both of these scores are indicative of the fact that they have the best feature selection. Further, these neural networks have the highest Model Accuracy (92%) of any neural network, which illustrates their ability to maintain prediction performance despite the presence of privacy limits.

When compared to more modern models, such as Logistic Regression and Decision Trees, traditional models have a larger information loss (18% and 15%, respectively) and lower k-anonymity levels (3 and 4), which implies that they are less effective at protecting the privacy of individuals. Support

Vector Machines (SVM) and Random Forests are two methods that perform relatively well in comparison to neural networks. These methods have accuracy rates of 89% and 90%, respectively. On the other hand, these methods are flawed when it comes to striking a balance between the dangers to privacy and the selection of features. All things considered, Privacy-Preserving Neural Networks are a fantastic choice for applications in sensitive areas since they display the best balance between privacy assurances and model functionality. This makes them an outstanding choice.

V. CONCLUSION

This paper aims to provide a thorough framework for machine learning models that protect privacy. Data anonymisation, privacy-aware feature selection, and privacy-protecting neural networks are combined to create this system. The presented methodology addresses significant challenges related to sensitive data protection while maintaining a high degree of model utility. Using techniques like k-anonymity and l-diversity, data is appropriately anonymised during the preprocessing stage. This guarantees minimum information loss while maintaining adequate privacy protection. High-utility, low-risk characteristics are preserved when privacy-conscious features are chosen, which maximises the feature set and reduces the quantity of sensitive data that is shown. Classification using privacy-preserving neural networks shows that it is feasible to attain competitive accuracy while also ensuring privacy by showcasing the usefulness of secure computation frameworks like homomorphic encryption and differential privacy. The evaluation metrics have demonstrated that the framework may effectively strike a balance between prediction performance and data privacy. This approach works especially effectively in highly sensitive sectors of the economy where ethical data handling is crucial, like healthcare and banking. This strategy could be improved upon in future studies by incorporating more advanced safeguards for user privacy and addressing issues with computing overhead.

REFERENCES

1. J. Abadi, A. Chu, I. Goodfellow, H. B. McMahan, I. Mironov, K. Talwar, and L. Zhang, "Deep learning with differential privacy," *Proceedings of the 2016 ACM SIGSAC Conference on Computer and Communications Security (CCS)*, Vienna, Austria, pp. 308–318, Oct. 2016. DOI: 10.1145/2976749.2978318.
2. R. Shokri and V. Shmatikov, "Privacy-preserving deep learning," *Proceedings of the 22nd ACM SIGSAC Conference on Computer and Communications Security (CCS)*, Denver, CO, USA, pp. 1310–1321, Oct. 2015. DOI: 10.1145/2810103.2813687.
3. Rani, S., Ghai, D., & Kumar, S. (2022). Reconstruction of Simple and Complex Three Dimensional Images Using Pattern Recognition Algorithm. *Journal of Information Technology Management*, 14(Special Issue: Security and Resource Management challenges for Internet of Things), 235-247. doi: 10.22059/jitm.2022.87475
4. P.S. Ranjit, V. Chintala, "Direct utilization of preheated deep fried oil in an indirect injection compression Ignition engine with waste heat recovery framework. *Energy*, 242, 122910, 2022. Elsevier (SCI), <https://doi.org/10.1016/j.energy.2021.122910>
5. William DeGroat, Dinesh Mendhe, AtharvaBhusari, HabibaAbdelhalim, SamanZeeshan, Zeeshan Ahmed, IntelliGenes: a novel machine learning pipeline for biomarker discovery and predictive analysis using multi-genomic profiles, *Bioinformatics*, Volume 39, Issue 12, December 2023, btad755, <https://doi.org/10.1093/bioinformatics/btad755>
6. Kumar, S., Raja, R., Mahmood, M.R. *et al.* A Hybrid Method for the Removal of RVIN Using Self Organizing Migration with Adaptive Dual Threshold Median Filter. *Sens Imaging* **24**, 9 (2023). <https://doi.org/10.1007/s11220-023-00414-9>
7. B. Dwork, C. McSherry, K. Nissim, and A. Smith, "Calibrating noise to sensitivity in private data analysis," *Theory of Cryptography Conference (TCC)*, Cambridge, MA, USA, pp. 265–284, Feb. 2006. DOI: 10.1007/11681878_14.
8. Kumar S, Choudhary S, Jain A, Singh K, Ahmadian A, Bajuri MY. Brain Tumor Classification Using Deep Neural Network and Transfer Learning. *Brain Topogr.* 2023 May;36(3):305-318. doi: 10.1007/s10548-023-00953-0. Epub 2023 Apr 15. PMID: 37061591.
9. R. Nuthakki, P. Masanta and Y. T N, "Speech Enhancement based on Deep Convolutional Neural Network," 2021 Fifth International Conference on I-SMAC (IoT in Social, Mobile, Analytics and Cloud) (I-SMAC), 2021, pp. 1-6, doi: 10.1109/I-SMAC52330.2021.9640736.
10. Rajakumari, R. F., & Ramkumar, M. S. (2021). Design considerations and performance analysis based on ripple factors and switching loss for converter techniques. *Materials Today: Proceedings*, 37, 2681-2686.
11. William DeGroat, Dinesh Mendhe, AtharvaBhusari, HabibaAbdelhalim, SamanZeeshan, Zeeshan Ahmed, IntelliGenes: a novel machine learning pipeline for biomarker discovery and predictive analysis using multi-genomic profiles, *Bioinformatics*, Volume 39, Issue 12, December 2023, btad755, <https://doi.org/10.1093/bioinformatics/btad755>
12. C. Gentry, "Fully homomorphic encryption using ideal lattices," *Proceedings of the 41st Annual ACM Symposium on Theory of Computing (STOC)*, Bethesda, MD, USA, pp. 169–178, May 2009. DOI: 10.1145/1536414.1536440.
13. P. S. Ranjit, et.al., Direct utilization of straight vegetable oil (SVO) fromSchleicheraOleosa (SO) in a diesel engine – a feasibility assessment, *International Journal of Ambient Energy*, DOI: 10.1080/01430750.2022.2068063
14. H. B. McMahan, E. Moore, D. Ramage, S. Hampson, and B. A. y Arcas, "Communication-efficient learning of deep networks from decentralized data," *Proceedings of the 20th International Conference on Artificial Intelligence and Statistics (AISTATS)*, Fort Lauderdale, FL, USA, pp. 1273–1282, Apr. 2017.
15. L. Zhu, Z. Liu, and S. Han, "Deep leakage from gradients," *Advances in Neural Information Processing Systems*

(*NeurIPS*), Vancouver, BC, Canada, vol. 32, pp. 14747–
14756, Dec. 2019.